

Notes on minimizing synthesized controller strategies

Andreas Holck Høeg-Petersen

September 2, 2022

1 Introduction

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetur id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

2 Overview

The process involves several steps.

The first step is to synthesize a strategy for a controller in UPPAAL Stratego and output it as a `json` file. UPPAAL stores the strategy as a set of Q-trees (one for each possible action). A Q-tree is a binary tree whose (internal) branching nodes each split the state space according to some bound on a given variable and whose leaf nodes are labelled with the Q-value of one particular action (the one that the given tree represents) in the state determined by the constraints on the path from the root node to that leaf. As such, the set of Q-trees together represents a complete lookup-table as known from classical RL, except that the discretization (partitioning) of the state space is unique for each action and is something that is learned during training instead of being pre-determined by experts or engineers.

The second step is to convert the set of Q-trees into an actual decision tree, where the leaf nodes are labelled not with Q-values but with actions. The decision tree should be constructed so that for each state S , the action label on the leaf node corresponding to S is always the same as the action with largest Q-value when evaluating S in all the original Q-trees. An algorithm for performing this conversion is given in Algorithm ???. In short, it works by obtaining all the state, action and Q-value tuples of the Q-trees and sorting them according to Q-value. The first tuple is then used to create a root node of the new DT and for each constraint needed to specify the state, a branching node is created. At the end of this path, a leaf node with the action is inserted. Now, the rest of the tuples can be inserted in order, creating branch nodes when the state needs to be further specified or discarding the insertion otherwise (since every new insertion will have a ‘worse’ Q-value than every previous one).

Thirdly, the resulting DT, that represents a single, complete partitioning of the state space with individual optimal actions assigned to every possible state, can then be scanned for ‘neighbouring’ leaves with identical action labels and aligned constraints. ‘Neighbouring’ in this sense refers to the situation where the states of two leaves are spatially adjacent and therefore could be represented by a single leaf instead of two. This will happen quite often, since UPPAAL will create a lot of arbitrary or experimental partitions especially during the early stages of training, which is then carried over into our new DT. This can —

potentially — give larger areas of the state space, where the same action is optimal but where multiple leaves (partitions) are used to represent it in the tree. Algorithm ?? describes a procedure for minimizing a given complete partitioning in the number of partitions used.

The fourth and final part is to recreate a DT from the newly obtained minimal partitioning. This has the challenge that the partitions most likely won't be able to be represented exactly by a tree structure. For example, the partitioning in Figure ?? cannot be represented by a binary tree, since no matter what (axis aligned) constraint we would choose for a root node, it would 'cut' one of the partitions in two and thus create a larger set of leaves than the set of partitions in the minimal partitioning. Instead, we then need a heuristic for how to best represent the minimal partitioning, which we give a suggestion for in Algorithm ??.

After this process, we can export the newly created DT to a format readable by UPPAAL Stratego and verify, that it performs just as well on the original task as the original strategy did.

3 Q-trees

In Reinforcement Learning (RL) the goal is to learn a policy π for which action to take in any given state of the environment. This policy is thus essentially a mapping from a state $S \in \mathcal{S}$ to an action $A \in \mathcal{A}$ where \mathcal{S} is the state space and \mathcal{A} is the set of all possible actions. Most often, the policy will be to choose the optimal action in any given state according to some metric, for example expected cumulative reward or cost. This is called a *greedy* policy.

The classical approaches either consider S to be a vector of continuous values and the task of learning π as one of approximating the function $f : S, A \mapsto Q_{S,A}^\pi$ where $Q_{S,A}^\pi$ is the expected value of taking action A in state S when following π (this can for example be done using a Neural Network ??), or considers S to describe a *discrete* state for which the true value of $Q_{S,A}^\pi$ can be learned for any $S \in \mathcal{S}$ and any $A \in \mathcal{A}$. In the latter case, if the state space originally is continuous, some kind of artificial discretization is required and then the strategy will typically be represented by a look-up table with dimensions $|\mathcal{S}| \times |\mathcal{A}|$. This is called a Q-table.

In UPPAAL Stratego ?? the gap between a continuous and discretized state space is gapped by having the discretization happen as part of the training process and representing the Q-values of state-action pairs as values in the leaf nodes of binary decision trees. During training, UPPAAL Stratego will build such trees by creating branch nodes that splits the state space in a particular dimension on a bound that it finds (or guesses) separates diverging Q-values of an action. These predicates take the form of $x \leq b$ where x is a dimension (variable) in the state space and b is a real valued constant.




Figure 1: A simple example of Q-trees for three actions, α , β and γ , and predicates on two variables, x and y . Evaluating a state with $x = 3$ and $y = 5$, we would get Q-values 10.0, 13.0 and 3.0 for α , β and γ respectively, meaning γ would be the optimal action.

The end result is a set of trees \mathcal{T} , one for each action $A \in \mathcal{A}$, where any continuous valued state $S \in \mathcal{S}$

What is the correct and concise way to write 'the dimensionality is the number of discrete states times the number of actions'?

can be evaluated to $Q_{S,A}^\pi$ in any tree $T_A \in \mathcal{T}$ by following the path from the root node to a leaf node as specified by the predicates in the branching nodes. That is, at each branch node we follow the path to the ‘left’ if the predicate is true for S and we follow the path to the ‘right’ otherwise, untill we end up at a leaf node which will then carry the value of $Q_{S,A}^\pi$. The subscript A in T_A indicates that it is the tree specifying the Q-values of action A and we can now define $T_A(S) : S, A \mapsto Q_{S,A}^\pi$. The greedy policy then just have to choose an action A where $A = \arg\min_{A \in \mathcal{A}} T_A(S)$ ¹.

Should I ditch the π superscript now?

We propose to call these trees *Q-trees* and provide a small example in Figure 1.

4 Converging Q-trees to decision tree

A strategy represented by a set of Q-trees has a couple of disadvantages: first of all, having to iterate all the trees and evaluate S in each one of them creates an amount of overhead, and second, the structure of the representation is not very intuitive to humans in the sense that the unique evaluation of S in each tree and the Q-values that they then return are very hard to interpret.

I don’t know

We can convert the Q-trees into a single decision tree with actions — not Q-values — in the leaf nodes so that the evaluation of S in the tree yields the optimal action for that state.

Maybe just cut this paragraph?

Let $\mathcal{V} = \{V_1, V_2, \dots, V_m\}$ be the set of all variables (dimensions) in the state space. We then have that a state $S = [v_1, v_2, \dots, v_m]^T$ is a vector of instantiated values of each $V \in \mathcal{V}$, ie. $V_i = v_i$. Following the path from the root node of any Q-tree T_A to a leaf node L_i , we can collect the constraints entailed by each predicate in the branching nodes along the path to get a symbolic state S_i that represents an m -dimensional area (also called partition) of the state space. The symbolic state is not given by concrete values of the dimensions but rather by lower and upper bounds, that is $S_i = \{(l_{1,i}, u_{1,i}), (l_{2,i}, u_{2,i}), \dots, (l_{m,i}, u_{m,i})\}$ where $l_{j,i}$ and $u_{j,i}$ are the lower and upper bounds respectively of V_j in the symbolic state S_i that the leaf L_i entails.

With this, we can define a leaf in a Q-tree as $L_i = (S_i, q_{A_i}, A_i)$ where S_i is the symbolic state obtained from applying all the constraints on the path to L_i , A_i is the action pertaining to the tree that the leaf belongs to and q_{A_i} is the Q-value of action A_i in all $S \in S_i$. The algorithm for converting a set of Q-trees to a decision tree then starts by gathering all these leaf triplets and sorting them in ascending order² according to their Q-value, and then pops the first leaf to create a root node and a path of branching nodes until S_1 is fully specified. Then a decision leaf can be inserted with label A_1 , knowing that this is the optimal action for this state. From here, we can continually pop and insert L_2, L_2, \dots, L_n in that order, creating new branching nodes when needed and ignoring insertions into areas that have already been covered by previous leaves (since the actions of these will necessarily be more optimal than the later one because of initial sorting).

5 Minimizing the decision tree

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetur id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

¹Note that UPPAAL Stratego uses *cost* rather than *reward*, therefore we minimize instead of maximizing

²Again, UPPAAL Stratego computes for Q-values the expected *cost*, so the action with the lowest Q-value is the most preferable.

5.1 Empirical pruning

Nam dui ligula, fringilla a, euismod sodales, sollicitudin vel, wisi. Morbi auctor lorem non justo. Nam lacus libero, pretium at, lobortis vitae, ultricies et, tellus. Donec aliquet, tortor sed accumsan bibendum, erat ligula aliquet magna, vitae ornare odio metus a mi. Morbi ac orci et nisl hendrerit mollis. Suspendisse ut massa. Cras nec ante. Pellentesque a nulla. Cum sociis natoque penatibus et magnis dis parturient montes, nascetur ridiculus mus. Aliquam tincidunt urna. Nulla ullamcorper vestibulum turpis. Pellentesque cursus luctus mauris.

5.2 Maximizing partition sizes

We start with a decision tree \mathcal{T} representing a complete partitioning of the state space. Let \mathcal{V} be the set of variables in the state space (ie. its dimensions) and let \mathcal{L} be the set of all leaves in the tree, where each leaf L_i has a state $S_i = \{(l_{i,1}, u_{i,1}), \dots, (l_{i,m}, u_{i,m})\}$ consisting of a list of tuples of lower and upper bounds for each of the m variables in \mathcal{V} . This state represents one m -dimensional box or area in the state space.

We then define \mathcal{C} as the set of constraints in the partitioning given by all the upper bounds, ie. $\mathcal{C} = \{u_{1,1}, \dots, u_{1,m}, u_{2,1}, \dots, u_{n,m}\}$. Lastly, we define a mutable list \mathcal{P} in which we store all the points from which our searches for mergable partitions begin. If we denote the minimum value of each variable $V_i \in \mathcal{V}$ as v_i^{\min} then we initialise \mathcal{P} as $\mathcal{P} = \{(v_1^{\min}, \dots, v_m^{\min})\}$.

The algorithm works by continually popping the next element of \mathcal{P} and using this point p as a starting point. From p we construct an infinitesimally small box $b = \{(p_1, p_1), \dots, (p_m, p_m)\}$ and then query \mathcal{T} to find out what action a_p is chosen at this symbolic state (which at this point is the same as querying at point p). Then we pick the constraint $u_{i,j} \in \mathcal{C}$ whose euclidean distance to the upper bound of b in dimension j is the smallest and update b to reflect this new upper bound in dimension i . That is, we replace $(p_j, p_j) \in b$ with $(p_j, u_{i,j})$.

Now we again query \mathcal{T} for the actions in b and continue updating b until we get more than one type of action. We then know, that the latest update of b expanded the box into an area of the state space where a different action is optimal, and thus we should not include it in our current box. We therefore roll back the latest update and mark that particular dimension as ‘exhausted’. If there are still unexhausted variables (dimensions) we continue as before, but now discarding any constraints on the exhausted dimensions.

If after we have rolled back our update, we see that all variables are exhausted, we then have a box b that spans the largest possible area (hyperplane?) from p with only one optimal action. We store this and then add to \mathcal{P} the edge points of b . That is, if $b = \{(l_1, u_1), \dots, (l_m, u_m)\}$ then we let $\mathcal{P} = \mathcal{P} \cup \{(u_1, l_2, \dots, l_m), (l_1, u_2, \dots, l_m), \dots, (l_1, l_2, \dots, u_m)\}$. After that, we start over with the next point in \mathcal{P} and continue until we have emptied \mathcal{P} .

Below is the pseudocode for the algorithm which includes a couple of more checks and edge cases needed for the algorithm to work. However, several details are still left out or relegated to (hopefully) helpfully named functions.

6 Rebuilding the tree

Fusce mauris. Vestibulum luctus nibh at lectus. Sed bibendum, nulla a faucibus semper, leo velit ultricies tellus, ac venenatis arcu wisi vel nisl. Vestibulum diam. Aliquam pellentesque, augue quis sagittis posuere, turpis lacus congue quam, in hendrerit risus eros eget felis. Maecenas eget erat in sapien mattis porttitor. Vestibulum porttitor. Nulla facilisi. Sed a turpis eu lacus commodo facilisis. Morbi fringilla, wisi in dignissim interdum, justo lectus sagittis dui, et vehicula libero dui cursus dui. Mauris tempor ligula sed lacus. Duis cursus enim ut augue. Cras ac magna. Cras nulla. Nulla egestas. Curabitur a leo. Quisque egestas wisi eget nunc. Nam feugiat lacus vel est. Curabitur consectetur.

Algorithm 1 MaxPartitions

```
1: function MAXPARTITIONS( $\mathcal{T}, \mathcal{C}, \mathcal{V}, p_{start}$ )
2:    $boxes \leftarrow \{\}$ 
3:    $\mathcal{P} \leftarrow \{p_{start}\}$ 

4:   while  $\mathcal{P}$  is not empty do
5:      $exhausted \leftarrow \{\}$ 
6:      $\mathcal{P}.sort()$  ▷ left to right
7:      $p \leftarrow \mathcal{P}.pop()$ 

8:     if  $p$  has been visited then ▷ can be checked by maintaining a separate tree
9:       continue

10:     $b \leftarrow \text{MAKEBOXFROMPOINT}(p)$ 
11:     $\text{SORTACCORDINGTODISTANCE}(\mathcal{C}, p)$ 

12:    for  $u_{i,j}$  in  $\mathcal{C}$  do
13:      if  $u_{i,j} \leq p_j$  or  $V_j \in exhausted$  then ▷  $p_j$  is value of  $V_j$  in  $p$ 
14:        continue

15:       $oldVal \leftarrow b_j^u$  ▷  $b_j^u$  is the upper bound of variable  $j$  in  $b$ 
16:       $b_j^u \leftarrow u_{i,j}$ 

17:      if  $\text{COUNT}(\mathcal{T}.actionsAtState(b)) > 1$  or  $b$  is explored or  $b_j^u = V_j^{\max}$  then
18:         $exhausted.add(V_j)$ 
19:        if  $exhausted.size < \mathcal{V}.size$  then
20:          continue

21:        if  $b_j^u < V_j^{\max}$  then ▷  $V_j^{\max}$  is the maximum possible value of  $V_j$ 
22:           $b_j^u \leftarrow oldVal$ 
23:           $boxes.add(b)$ 

24:        for  $p'$  in  $\text{EDGEPOINTS}(b)$  do
25:           $\mathcal{P}.add(p')$ 
26:  return  $boxes$ 
```

7 Experiments

7.1 Maximizing partitions

Continuing with the bouncing ball example, we now try and apply the **MaxPartitions** algorithm (see Section 5.2) to the strategy trained by UPPAAL Stratego. This first step is, as mentioned, to convert the set of Q-trees into a decision tree. The reason for this is that it is one way to merge the (two) action specific partitionings of the state space into one single partitioning where every area is prescribed one optimal action.

Due to the randomness involved in how the decision tree turns out (caused by the fact that we randomize the order in which state bounds are turned into branching nodes), we cannot expect exactly the same tree size at every instance of performing the conversion. For this experiment, we did the conversion 100 times and only saved the best (smallest) version of the tree. This gave us an instance of the complete strategy represented by a tree with 86,807 paths from the root to a leaf.

Since the bouncing ball example only has two state variables (position and velocity) we can actually draw the state partitioning and colorize the different partitions according to the action that is prescribed (green for no hit, red for hit). This can be seen in Figure 2.




Figure 2: A 2D visualization of the partitioning (dotted lines) of the state space of the bouncing ball model as it looks without any minimization attempts. The x-axis is velocity and the y-axis is the balls position. Green areas represents the action ‘no hit’, red areas represents ‘hit’. The thick black line is $x = 0$ and there have been drawn grid lines at every 1 tick.

An inspection of this partitioning shows that a lot of partitions that are next to each other share the same action which make them fit for combining, thus reducing the overall number of partitions and maximizing the size of the partitions. Using the **MaxPartitions** algorithm we obtain a new set of partitions (not a DT yet) but the number has been dramatically reduced to only 703, that is, a 99.19% reduction! The new partitioning can be seen in Figure 3.

This new partitioning describes exactly the same strategy as the original Q-tree based representation from UPPAAL and the converted DT, but with significantly fewer individual parts. Converting it into a new decision tree gives a tree with 895 paths, that is, the conversion could not perfectly preserve the minimality of the new partitioning. Using this DT as an input strategy for the UPPAAL model, we see that it performs equivalently to the original UPPAAL Strategy (see Table 1).

Next, we perform empirical pruning (see Section 5.1) on the DT after running **MaxPartitions**. The sample was generated by running 1000 simulations of 300 timesteps each, logging the state at every 0.5 timestep. Pruning unvisited states then reduced the number of paths to just 189 while retaining performance. The tree was then further reduced by pruning redundant branch nodes and the final version of the strategy then only has 164 paths, but still manages to only require 38.311 hits in the evaluation test, just as the original UPPAAL strategy.



Figure 3: A visualization of the same strategy as in Figure 2 but now with much fewer and larger partitions.

Table 1: Comparing the performance of controllers for the bouncing ball example over 1000 runs for 120 timesteps each before and after various attempts at minimizing the size through either empirical pruning or the **MaxPartitions** algorithm.

| Version | Paths | Expectation (hits) | Deviation |
|--------------------------------------|--------|--------------------|-----------|
| Q-trees | 13,405 | 38.401 | 0.177 |
| Original DT | 86,807 | 38.431 | 0.178 |
| ZeroPruned | 6,256 | 39.900 | 0.205 |
| MaxPartitions | 895 | 38.435 | 0.177 |
| MaxPartitions then Prune | 189 | 38.347 | 0.173 |
| MaxPartitions, EmpPrune and RedPrune | 164 | 38.311 | 0.172 |

7.2 dtControl

The tool **dtControl** has the ability to take a synthesized strategy represented as a look-up table and convert the representation to a decision tree that respects the safety requirements but compresses the size immensely. This is naturally of great interest for our case, but both the Q-tree representation and our own decision tree conversion alters the setup somewhat.

However, even though the tool actually supports directly working with the output format of UPPAAL Stratego, this is only the case for strategies learned using the **control[]** directive, which controls for certain defined safety parameters to always be respected. In our case of the bouncing ball example, the controller is trained with the **minE[]** directive, which minimizes a given parameter (in this case, the number of times the controller hits the ball).

Instead, we had to create our own output files to use as input for **dtControl**. According to the doc-

umentation, a controller strategy can be specified in a simple CSV format where each line contains an allowed state/action pair, that is, N values representing a state where the following M values constitutes an allowed action. For example, in the case of the bouncing ball, we have two state variables (position and velocity) and one action variable (hit or not hit), meaning each line would have three values.

We have attempted two experiments with different ways of specifying our original controller.

In the first experiment, we used the trained controller to generate 30,000 samples of state/action pairs. That is, the UPPAAL model was run for 300 timesteps with the trained controller deciding what action to take in each encountered state and then the state/action pairs were logged at each 0.01 timestep.

In the second experiment, we converted the strategy a set of Q-trees (the initial UPPAAL format) to a DT as described in Section 4. This DT had a partition size (number of leaves/paths) of 91,054. We used these partitions as the input data to `dtControl` by taking the maximum value of each variable in each individual partition together with the optimal action of that state. That is, we effectively specified the discretization of the state space by defining the bounds of each state paired with the allowed/optimal action.

Table 2: Comparing the performance of controllers for the bouncing ball example over 1000 runs for 120 timesteps each before and after various attempts at minimizing the size through `dtControl`.

| Version | Paths | Construction time | Expectation (hits) |
|--------------|--------|-------------------|--------------------|
| Original DT | 91,054 | — | 38.431 |
| Samples | 27,234 | 8:14 | 318.411 |
| State bounds | 521 | 0:41 | 315.769 |

The results when applying the generated strategies to the model in UPPAAL are given in Table 2 together with the baseline original decision tree directly converted from the Q-tree set. As is seen, when we used samples, we still got a somewhat large DT that took more than 8 minutes to generate. And the performance (expected number of hits) is substantially worse than the baseline version. For the version based on state bounds, we got a much smaller tree with only 521 paths, but the performance was still very far from the original.